

CS 503 Fall 2022: Lab 1

Suraj Aralihalli

September 2022

3. Inspecting and modifying XINU source, compiling, and executing on backend machines

3.1 Basic system initialization and process creation

1. **Question:** Why the output is consistent with correct operation of `create2()`?

Solution: The output i.e sequence of A's and B's is consistent and matches the output of `create()` because in my implementation of `create2()` I use the function `ready()` to change the state, add the process to the ready queue and reschedule. Additionally, I am using the same priority for all the processes. Each process runs for Quantum unit of time before its switched. Process "send A" prints a sequence of As for 1 quantum before it is context switched to process "send B" which prints a sequence of Bs. Since I have used the same priority value in `create()` and `create2()` while creating the processes and in the same order ($A \rightarrow B$), the output printed on the screen is consistent w.r.t both `create()` and `create2()`.

2. **Question:** Trace the sequence of events that transpire when the function executed by `create()` returns.

Solution: When the function executed by `create()` returns, the control flows to the instructions at `INITRET`. `INITRET` is an alias to `userret`. Function `userret()` is responsible for killing the function from its PID. `userret()` function uses `kill()` system call which changes the state of the process to `PR_FREE` and calls `resched()` to maintain fixed priority scheduling. `kill()` also sends a message to the parent process. To reiterate the sequence of steps undertaken,

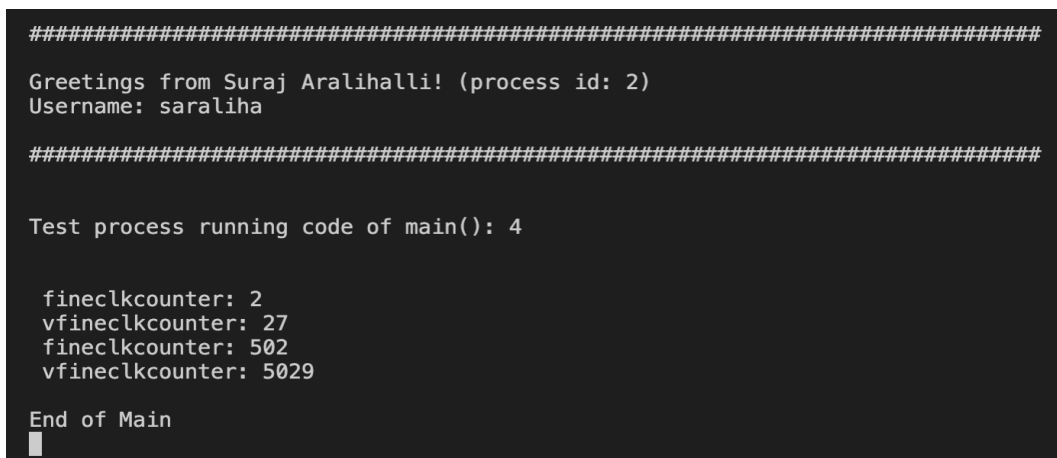
- (a) Control flows to `userret()`
- (b) `userret()` captures `cur pid` and calls `kill(pid)`
- (c) `kill` saves the interrupt mask and performs sanity checks
- (d) Decrements the `prcount` (process count)
- (e) Sends a message to the parent process
- (f) Frees the stack space of the process created by `create()` by calling `freestk()`
- (g) Changes the state of the child process id to `PR_FREE`
- (h) Calls `resched()` to maintain fixed priority scheduling
- (i) Restores the interrupt mask

3.2 Clock interrupt handling

Question: Test that `fineclkcounter` and `vfineclkcounter` are being updated correctly by XINU's modified clock interrupt handling code by making `main()` sleep for 5 seconds by calling `sleep()`, then printing the value of the two global variables.

Solution: See Figure 1

```
kprintf("\n fineclkcounter: %d", fineclkcounter);
kprintf("\n vfineclkcounter: %d", vfineclkcounter);
sleep(5);
kprintf("\n fineclkcounter: %d", fineclkcounter);
kprintf("\n vfineclkcounter: %d", vfineclkcounter);
```



```
#####
Greetings from Suraj Aralihalli! (process id: 2)
Username: saraliha
#####

Test process running code of main(): 4

fineclkcounter: 2
vfineclkcounter: 27
fineclkcounter: 502
vfineclkcounter: 5029

End of Main
█
```

Figure 1: `fineclkcounter` and `vfineclkcounter`

3.3 Time slice management

Question: Test whether `prepuhungry` accurately tallies how CPU intensive a process is by creating a test case.

Solution: My test consists of running two different processes and comparing the corresponding `prepuhungry` value. As seen below I have a function named `cpuIntensiveLow()` that runs 10 million times and computes $sum = (sum + i - 1)\%i$; in each iteration. Similarly, I have another function named `cpuIntensiveHigh()` that runs the same operations but runs 100 million times. I then created the 2 processes in the main using `create()/create2()` and printed the `prepuhungry` value of the process right before the process returns. Since this value represents the number of quantum a process consumed from the time it was created to the time the process was killed, I would expect `cpuIntensiveHigh()` to run for more quantum than `cpuIntensiveLow()`. From the figure 3, it is evident that `prepuhungry` accurately tallies how CPU intensive a process as per my tests.

```
void cpuIntensiveLow()
{
    int i = 0;
```

```

        int sum = 0;
        while(i!=10000000)
        {
            sum=(sum+i-1)%i;
            i++;
        }
    }

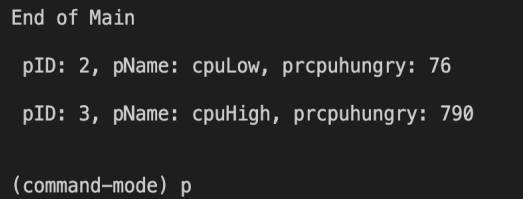
void cpuIntensiveHigh()
{
    int i = 0;
    int sum = 0;
    while(i!=100000000)
    {
        sum=(sum+i-1)%i;
        i++;
    }
}

process main(void)
{
    kprintf("\nTest process running code of main(): %d\n\n", getpid());

    create2(cpuIntensiveLow, 1024, 20, "cpuLow", 0);
    create2(cpuIntensiveHigh, 1024, 20, "cpuHigh", 0);

    kprintf("\n\nEnd of Main\n");
    return OK;
}

```



```

End of Main

pID: 2, pName: cpuLow, prcpuhungry: 76
pID: 3, pName: cpuHigh, prcpuhungry: 790

(command-mode) p

```

Figure 2: cpuIntensiveLow() and cpuIntensiveHigh()

3.4 Process lifetime

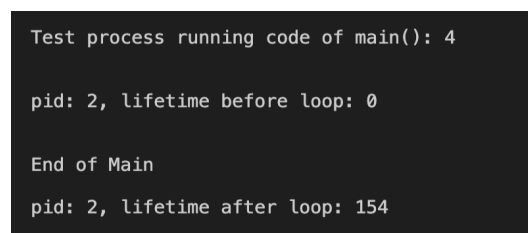
Question: Test and verify that the system call works correctly.

Solution: I modified the cpuIntensiveLow() function to display the lifetime() once before the for loop and once after the for loop. From the Figure ?? it is clear that the lifetime() accurately computes the total time the process was alive in the unit of 10ms.

```

void cpuIntensiveLow()
{
    kprintf("\npid: %d, lifetime before loop: %d\n", getpid(), lifetime(getpid()));
    int i = 1;
    int sum = 0;
    while(i!=10000000)
    {
        sum=(sum+i-1)%i;
        i++;
    }
    kprintf("\npid: %d, lifetime after loop: %d\n", getpid(), lifetime(getpid()));
}

```



```

Test process running code of main(): 4

pid: 2, lifetime before loop: 0

End of Main

pid: 2, lifetime after loop: 154

```

Figure 3: cpuIntensiveLow() with lifetime()

3.5 Energy conservation

Question: Use the tests of 3.3 and 3.4 to gauge that the energy efficient idle process appears to be working correctly.

Solution: Running the tests 3.3 and 3.4 after adding `asm("hlt")` gave similar results.

4. Interfacing C and assembly code

4.1 Calling assembly function from C function

Question: Test and verify that `addtwo()` and `addfour()` work correctly.

Solution: Functions `testAddtwo()` and `testAddfour()` can be found in `mytests.c` in `systems` directory. These functions are called in `main()` to test `addtwo()` and `addfour()` functions written in assembly. See Figure 4 for reference.

4.2 Calling C function from assembly function

Question: Test and verify that your code for `testgreaterfirst` and `greaterfirst` works correctly. **Solution:** `testgreaterfirst()` successfully accesses the arguments passed by its caller and communicates it to `greaterfirst`. It also ensures value contained in `EAX` is not disturbed and is returned to its caller. Function `testTestgreaterfirst()` used for testing can be found in `mytests.c` in `systems` directory.

```

.text
.globl testgreaterfirst

```

```

Test process running code of main(): 4

PASS: Adding 2 positive numbers 3+4=7
PASS: Adding 2 positive, negative numbers 3-4=-1
PASS: Adding 4 positive numbers 3+4+5+6=18
PASS: Adding 4 positive, negative numbers -3+4-5+6=2

End of Main

```

Figure 4: testAddtwo() with testAddfour()

```

testgreaterfirst:
    pushl    %ebp
    movl    %esp,%ebp
    pushfl
    pushl    %ebx
    movl    8(%ebp),%eax
    movl    12(%ebp),%ebx
    pushl    %ebx
    pushl    %eax
    call    greaterfirst
    popl     %ebx
    popl     %ebx
    popl     %ebx
    popfl
    movl     (%esp),%ebp
    add      $4,%esp
    ret

```

5. Run-time manipulation of return addresses

5.1 Wrong turn

Figure 5 shows the system entering panic state when greaterfirst1 is called from testgreaterfirst. Code snippet from greaterfirst1.c responsible for the trap follows.

```

asm("movl %ebp,ebp");
// ebp + 4 has return address
returnAddressToCaller = ebp + 0x1;
*returnAddressToCaller = 0x123456;

```

5.2 Expressway to main()

In greaterfirst2.c, ebp points to address that stores frame pointer to testgreaterfirst(). We then make ebpInCaller point to this address. We know that the return address to testgreaterfirst() is stored in the address pointed by ebp + 4bytes. Similarly return address to main() is stored in ebpInCaller + 4bytes. We now replace the return address to testgreaterfirst() with return address to main(). This

```

Test process running code of main(): 4

Xinu trap!
exception 13 (general protection violation) curripid 4 (Main process)
error code 00000000 (0)
CS EFC0008 eip 101EF5
eflags 10297
register dump:
eax 00000000 (0)
ecx FFFFFFFF (4294967295)
edx 0011698F (1141135)
ebx 00000002 (2)
esp 0EFC8F80 (251432832)
ebp 0EFC8F80 (251432832)
esi 00000000 (0)
edi 00000000 (0)

panic: Trap processing complete...

```

Figure 5: trapping when greaterfirst1 is called from testgreaterfirst

way when the machine instruction **call greaterfirst2** is completed, the control flows to main(). Below is the snippet from testgreaterfirst.S

```

unsigned long *ebpInCaller=NULL;
unsigned long *returnAddressToCaller=NULL;
unsigned long *returnAddressToMain=NULL;

asm("movl %ebp,ebp");

ebpInCaller = (unsigned long *)*ebp;

// ebp + 4 has return address
returnAddressToCaller = ebp + 0x1;
returnAddressToMain = ebpInCaller + 0x1;

*returnAddressToCaller = *returnAddressToMain;

```

6. Modifying interrupt handling behavior

```

process main(void)
{
    kprintf("\nBefore divide by zero \n");
    int o = 100;
    o = 5 / 0;
    kprintf("\nValue of k:%d \n", o);
    kprintf("\nAfter divide by zero \n");
    return OK;
}

```

The above code is used to test system's response to divides by zero operation. When Version1 of `_Xint0` is used, the system gets stuck in an infinite loop. This is confirmed from the Figure 6. The print statements immediately after `o = 5/0` are not printed, indicating that the process is in an infinite loop stuck at `o = 5/0` statement.

When Version2 of `_Xint0` is used, the process returns to the next statement after `o = 5/0`. Consequently, the print statements are executed and can be seen in Figure 6. This indicates that the process makes progress unlike the previous version and doesn't get stuck. In Version2 we know that the return address is at the top of the stack, we store this address into `ecx` register and add offset of 2 bytes to `ecx` (size of `idiv` instruction, See line 2c below). And we put it back on top of the stack without modifying anything else.

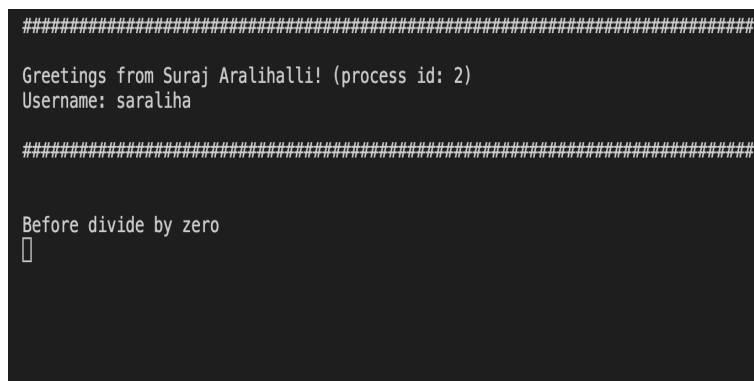
```
29:  c1 fa 1f          sar    $0x1f,%edx
2c:  f7 f9            idiv   %ecx
2e:  89 44 24 1c       mov    %eax,0x1c(%esp)
```

Version1:

```
.globl _Xint0
_Xint0:
    iret
```

Version2:

```
.globl _Xint0
_Xint0:
    popl %ecx
    addl $2,%ecx
    pushl %ecx
    iret
```



```
#####

Greetings from Suraj Aralihalli! (process id: 2)
Username: saraliha

#####

Before divide by zero
█
```

Figure 6: Version1 `_Xint0`: stuck in infinite loop

```
#####
Greetings from Suraj Aralihalli! (process id: 2)
Username: saraliha
#####

Before divide by zero

Value of k:100

After divide by zero
```

Figure 7: Version2 _Xint0: Process proceeds to next statement

Bonus problem

Code:

```
int createapp(char *cmd, char **argv)
{
    int pid = fork();
    //Child
    if (pid==0)
    {
        nice(10);
        // sleep(10);
        execvp(cmd, argv);
        // comes here only if execvp fails
        exit(1);
    }
    //Parent
    else
    {
        int status;
        // blocking call
        waitpid(pid, &status, 0);

        if ( WIFEXITED(status) )
        {
            int exit_status = WEXITSTATUS(status);
            if(exit_status==1)
            {
                return -1;
            }
        }
    }
    return pid;
}
```



```
}
```

My approach to fix the flaw where `createapp()` would return the PID of the child even if `execvp()` in the child fails, is to use a blocking system call from the parent that waits until the child executes successfully (exit status 0) or terminates (exit status 1). Based on this exit status, the parent

1. returns **pid** of the child if `execvp()` in child succeeds
2. returns **-1** if `execvp()` in the child fails

The same behaviour is also found in `system()` in Linux which returns only after the command has been completed. I have included the code snippet below where I use `waitpid()` with third argument 0. This makes the `waitpid()` a blocking system call and it returns only after the child is terminated. The status (success/failure) of the child process is captured in the `exit_status` variable. If the `exit_status` is 1, then we return -1. Otherwise, we return pid of the child as instructed in the question.

I tested the successful execution of `execvp()` by passing the valid arguments to `createapp()` function. On the other hand the I simulated the failure execution of `execvp()` by passing invalid arguments to `createapp()`. For example, in the first test case I passed valid arguments (See Testcase 1 below). I could see the pid of child process printed on the terminal as expected. When I passed invalid arguments in testcase 2 to intentionally fail the `execvp()` command, the parent process waits until the completion of the child process and identifies that the `exit_status` of child is 1 and prints -1 on terminal.

Testcase 1:

```
int main()
{
    char* command = "ls";
    char* argument_list[] = {"ls", "-l", NULL};
    int i = createapp(command, argument_list);
    printf("\npid: %d\n", i);
}
```

Testcase 2:

```
int main()
{
    char* command = "qdhshjdsaj";
    char* argument_list[] = {"qdhshjdsaj", "-l", NULL};
    int i = createapp(command, argument_list);
    printf("\npid: %d\n", i);
}
```